



for multiple regression

<https://andyigg.com/spa3-1/>

# Software

- R: <http://cran.r-project.org/>
- RStudio:  
<https://www.rstudio.com/products/rstudio/download/#download>

## Assessment of residuals

- Outliers
- Influential observations
- Normality of residuals
- Equal variance of residuals
- Independence of residuals

# Multiple Linear Regression Example

```
set.seed(108) ; n= 100 ; beta0 = array( c(3,2,1,1), c(1,4) )
```

```
x1=rgamma( n, 1, 1/10 ) ; x2=rchisq( n, df = 3 ) ; x3=rexp( n )
```

```
X= cbind( x1, x2, x3 )
```

```
y= beta0[,1] + beta0[,2]*x1 + beta0[,3]*x2 + beta0[,4]*x1*x2 +  
rnorm( n,sd=1.5 )
```

- `summary(mlm)`

```
Call:
lm(formula = y ~ x1 + x2 + x3 + x1:x2 + x1:x3)

Residuals:
    Min       1Q   Median       3Q      Max
-3.7979 -1.0429 -0.0360  0.9585  3.8980

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.770501   0.392648   7.056 2.89e-10 ***
x1           2.044134   0.031253  65.406 < 2e-16 ***
x2           1.075617   0.092715  11.601 < 2e-16 ***
x3          -0.225363   0.229452  -0.982  0.329
x1:x2        0.991471   0.008786 112.846 < 2e-16 ***
x1:x3       -0.003571   0.024526  -0.146  0.885
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.52 on 94 degrees of freedom
Multiple R-squared:  0.9989,    Adjusted R-squared:  0.9989
F-statistic: 1.731e+04 on 5 and 94 DF,  p-value: < 2.2e-16
```

- `confint(mlm, level=0.95)`


```
> confint(mlm, level=0.95)
                2.5 %      97.5 %
(Intercept)  1.9908882 3.55011356
x1           1.9820801 2.10618759
x2           0.8915291 1.25970485
x3          -0.6809452 0.23021915
x1:x2        0.9740257 1.00891550
x1:x3       -0.0522689 0.04512595
```

- `anova(mlm)`

```
> anova(mlm)
Analysis of Variance Table

Response: y
      Df Sum Sq Mean Sq    F value Pr(>F)
x1      1 132044  132044 57176.0092 <2e-16 ***
x2      1  36281   36281 15709.9284 <2e-16 ***
x3      1      0      0    0.1101 0.7408
x1:x2    1  31555   31555 13663.6479 <2e-16 ***
x1:x3    1      0      0    0.0212 0.8845
Residuals 94    217      2
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> |
```

# Outliers?



- `#install.packages("car") ; library(car)`  
`influencePlot(mlm)`

# Assessing outliers

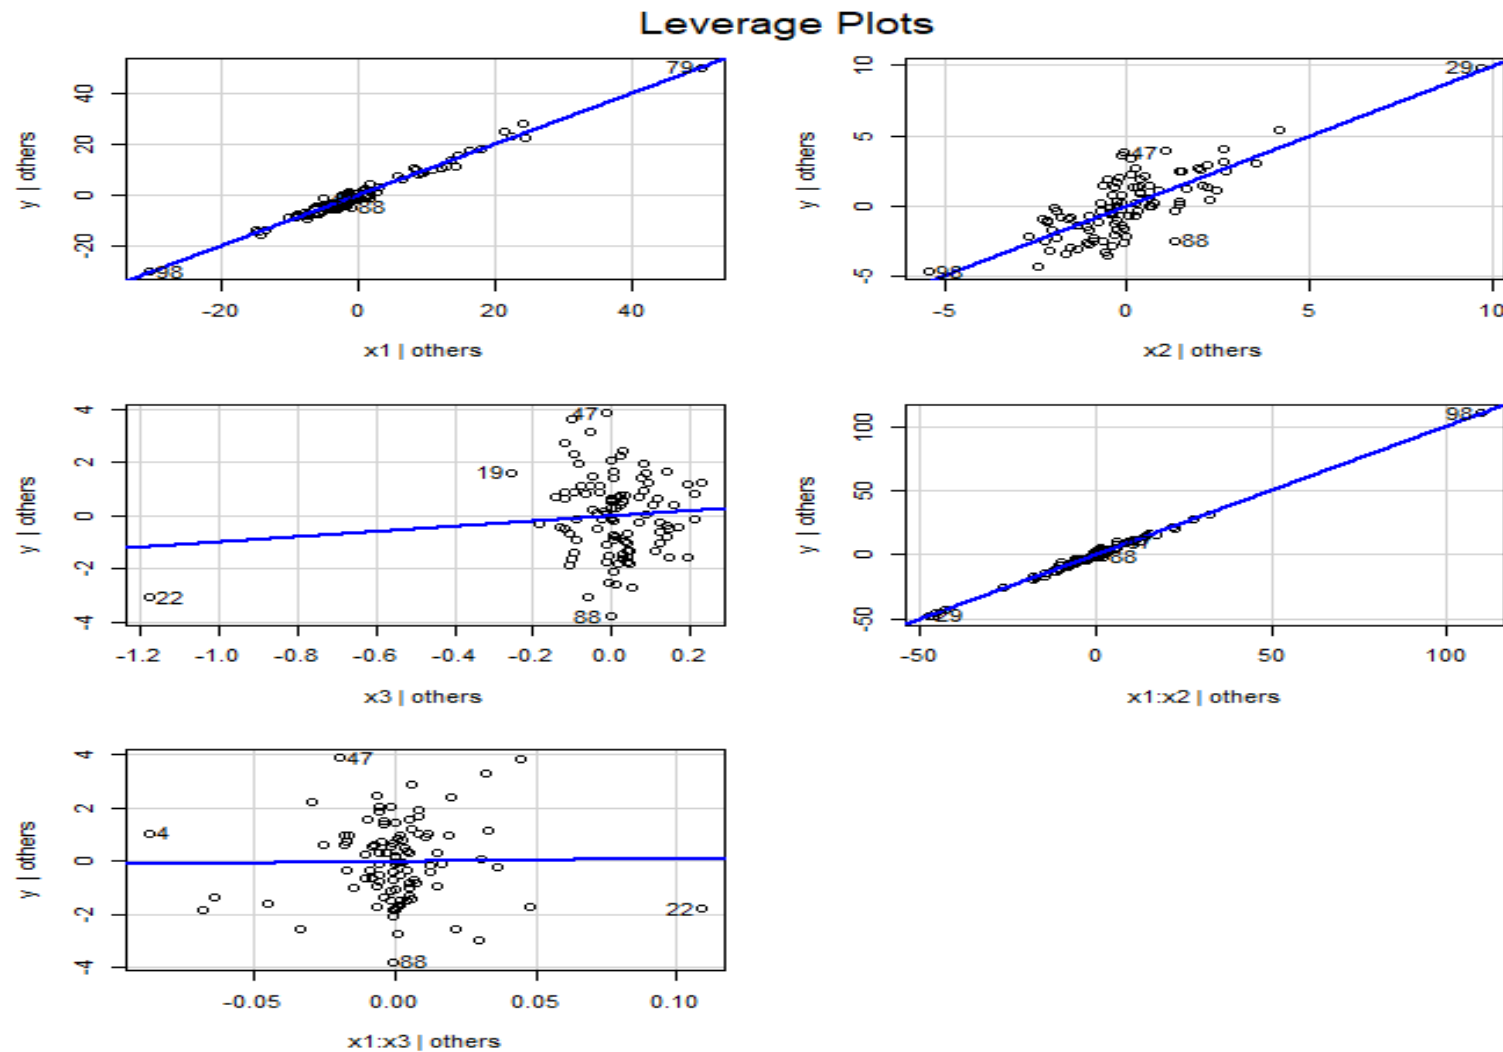
- outlierTest(mlm) # Bonferonni p-value

```
> outlierTest(mlm)
No Studentized residuals with Bonferroni p < 0.05
Largest |rstudent|:
      rstudent unadjusted p-value Bonferroni p
44 2.696387      0.0083212      0.83212
. 1
```

Now, Bonferroni:  $1 - \frac{\alpha}{k} = \alpha'$ , where k is the number of test


e.g.  $1 - (1 - \text{pnorm}(-2.696387))/6 = 1 - \text{pnorm}(2.696387)/6 \approx 0.832$

# leveragePlots(mlm) # leverage plots

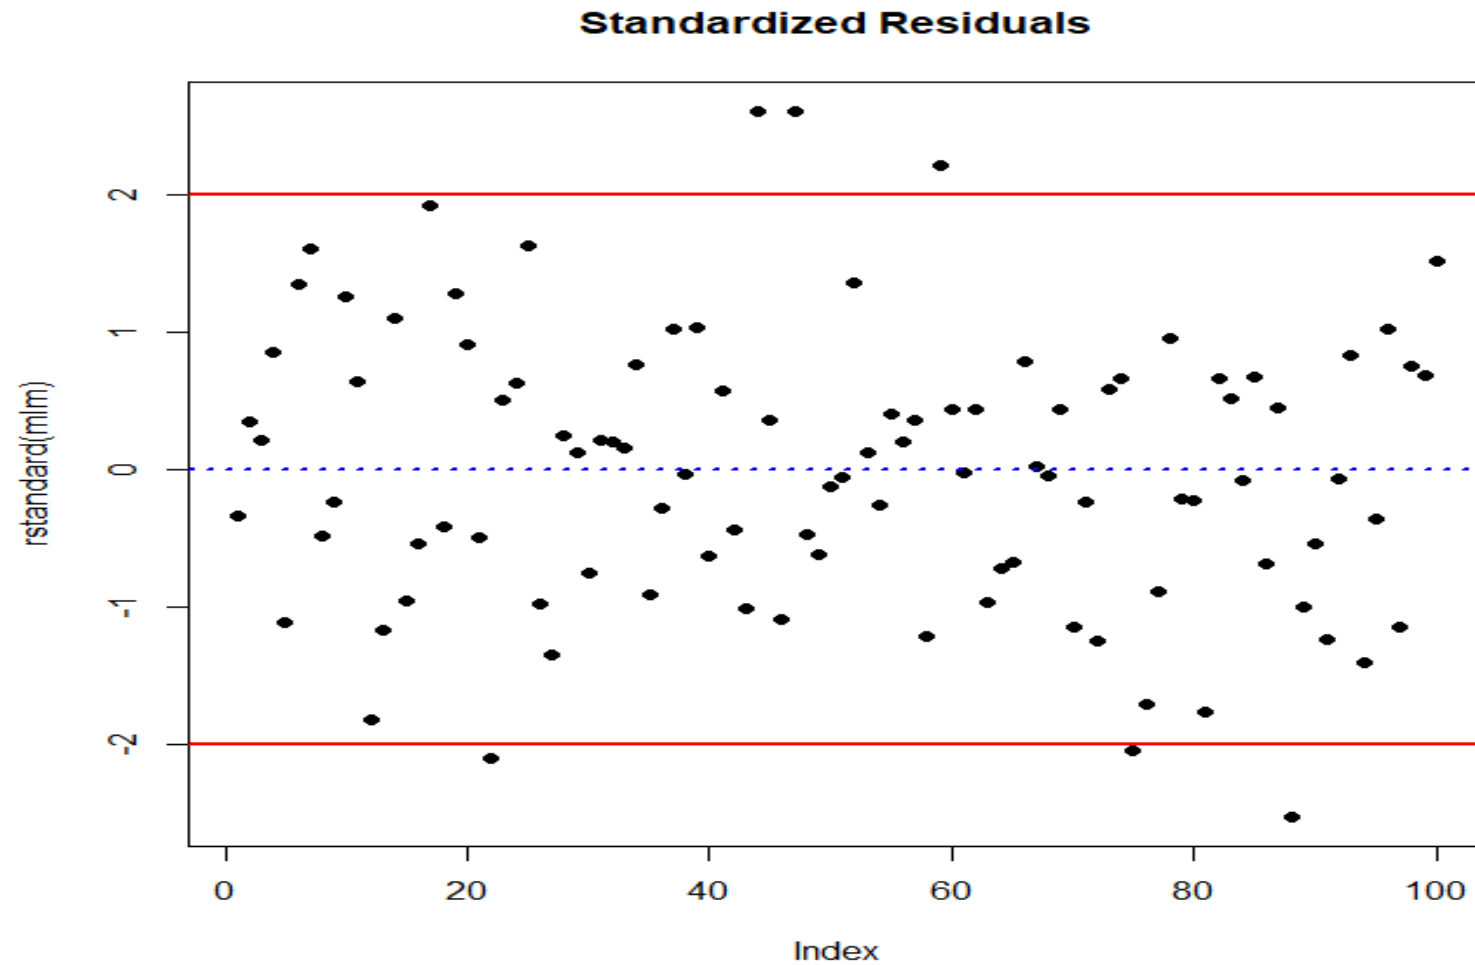




`plot(hatvalues(mlm))` # leverage  
plots



# Standardized residuals plot



```
plot(rstandard(mlm), type = "p", main =  
"Standardized Residuals")
```

```
points( 1:n, rstandard(mlm), pch = 19)
```

```
abline( a = -2, b = 0, lty = 1, lwd = 2, col = "red" )
```

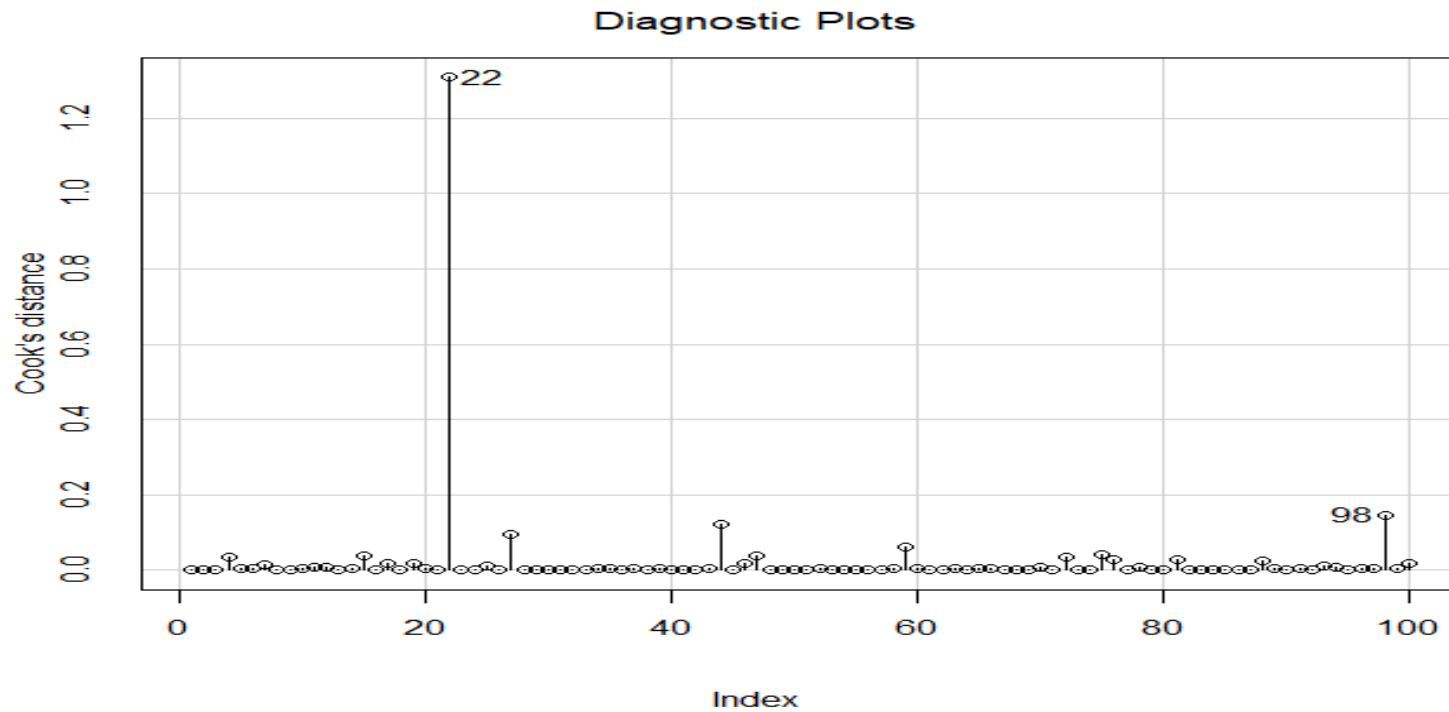
```
abline( a = 2, b = 0, lty = 1, lwd = 2, col = "red" )
```

```
abline( a = 0, b = 0, lty = 3, lwd = 2, col = "blue" )
```

# Influential observations: cook's distance

```
# Cook's D plot
# Identify D values > 4/(n-p-1)
cooks=cooks.distance(mlm)
cooks_mlm=cooks[cooks > 4/( n - length(mlm$coefficients ) )]
cooks_mlm

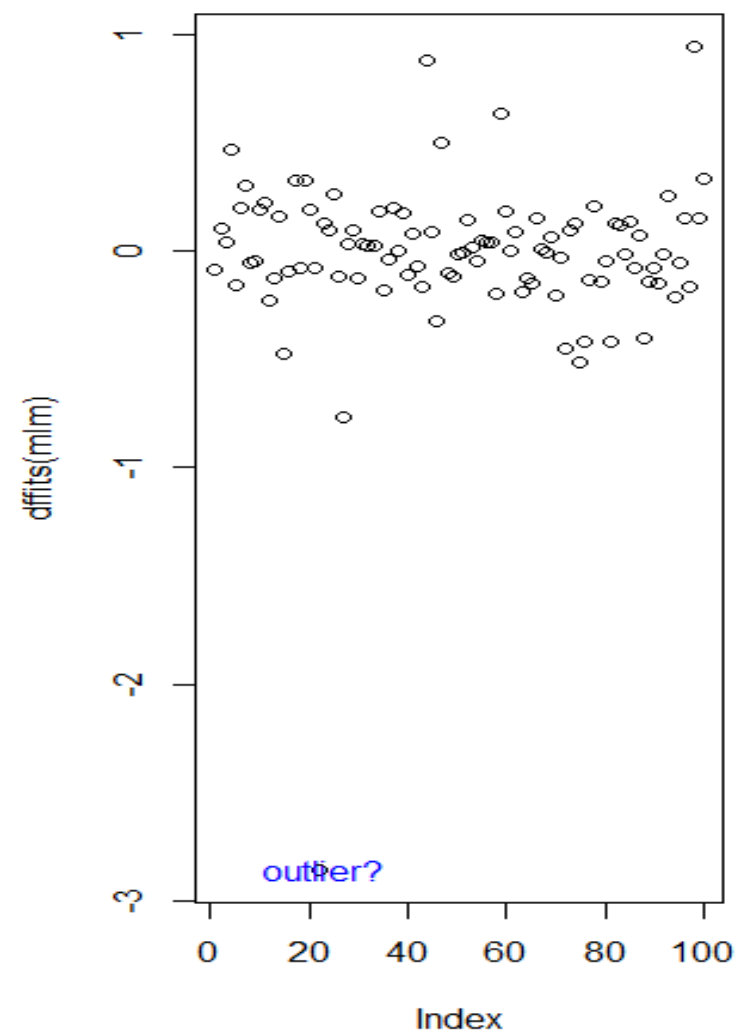
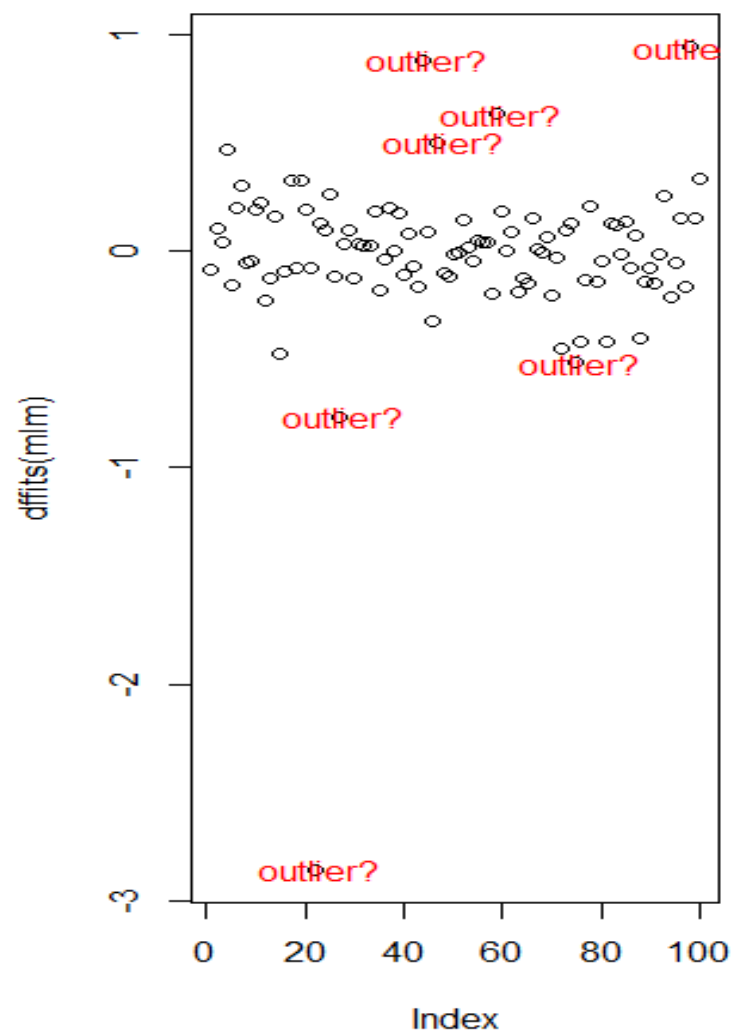
influenceIndexPlot(mlm, vars="Cook", id.n=length(cooks_out))
```



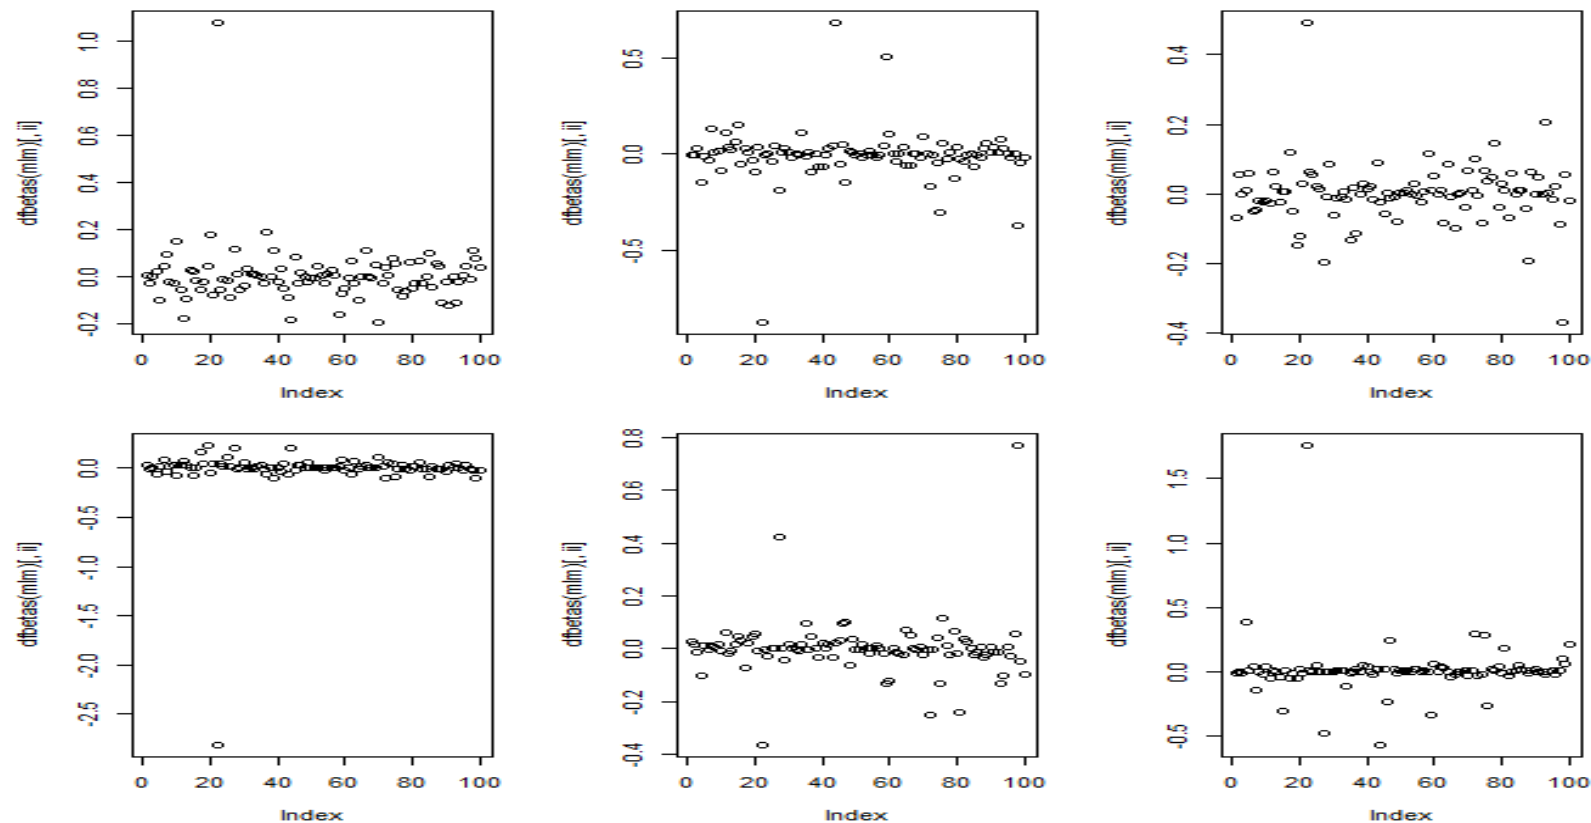
```
> cooks_mlm=cooks[cooks > 4/( n - length(mlm$coefficients) )]
> cooks_mlm
      22      27      44      59      75      98
1.30878691 0.09684502 0.12119111 0.06350814 0.04309227 0.14736014
>
```

# Influential observations: dffits

```
# dffits
par( mfrow=c(1,2), oma = c( 0, 0, 2, 0) ) ; par(mar=c(4,4,2,2))
plot(dffits(mlm)) ; # now, p=5
# dffits > sqrt((p+1)/n)
text( x=1:length(dffits(mlm))+1, y=dffits(mlm),
labels=ifelse(abs(dffits(mlm))>2*sqrt((5+1)/n), "outlier?", ""),
col="red" )
# dffits > 1)
plot(dffits(mlm)) ; # now, p=5
text( x=1:length(dffits(mlm))+1, y=dffits(mlm),
labels=ifelse(abs(dffits(mlm))>1, "outlier?", ""), col="blue" )
```





# Influential observations: dfbetas



```
par( mfrow=c(2,3), oma = c( 0, 0, 2, 0) ); par(mar=c(4,4,2,2))  
for( ii in 1:length(mlm$coefficients) ){ plot(dfbetas(mlm)[,ii]) }
```



# Normality of residuals



```
# library(MASS)
mlmresid=studres(mlm)
par( mfrow=c(1,2), oma = c( 0, 0, 2, 0) ) ;
par(mar=c(4,4,2,2))
qqnorm(mlmresid, ylab ="Standardized Residuals")
qqline(mlmresid, lty=1, lwd=2,col="red" )

qqPlot(mlmresid, main="QQ Plot") # QQ plot for
studentized residual
```

# Formal Tests for normality

- **shapiro.test(mlmresid)**
- **ks.test(mlmresid, "pnorm",  
mean(mlmresid), sd(mlmresid) )**
- **#require(nortest)**  
**lillie.test(mlmresid)**
- ...

```
> shapiro.test(mlmresid) # Shapiro-Wilks normality test
```

```
Shapiro-Wilk normality test
```

```
data:  mlmresid
```

```
W = 0.99349, p-value = 0.9155
```

```
> ks.test(mlmresid, "pnorm", mean(mlmresid), sd(mlmresid) ) # Kolmogorov-Smirnov normality test
```

```
One-sample Kolmogorov-Smirnov test
```

```
data:  mlmresid
```

```
D = 0.035714, p-value = 0.9996
```

```
alternative hypothesis: two-sided
```

```
> require(nortest)
```

```
> lillie.test(mlmresid) # Lilliefors (Kolomorov-Smirnov) normality test
```

```
Lilliefors (Kolmogorov-Smirnov) normality test
```

```
data:  mlmresid
```

```
D = 0.035714, p-value = 0.9895
```

# Which Normality Test Should I Use?

- **Kolmogorov-Smirnov test**

It is more sensitive near the center of the density than at the tails than other tests; large n

- **Shapiro-Wilks test**

Doesn't work well if several values in the data set are the same. Works best for data sets with small n, but can be used with larger data sets.

- **Lillie test** (adjusted Kolmogorov-Smirnov test)

More powerful than Kolmogorov-Smirnov test

[http://www.hmwu.idv.tw/web/R\\_AI/v2/hmwu\\_StatR-05-2\\_NonParametric\\_adv.pdf](http://www.hmwu.idv.tw/web/R_AI/v2/hmwu_StatR-05-2_NonParametric_adv.pdf)

<https://www.nrc.gov/docs/ML1714/ML17143A100.pdf>

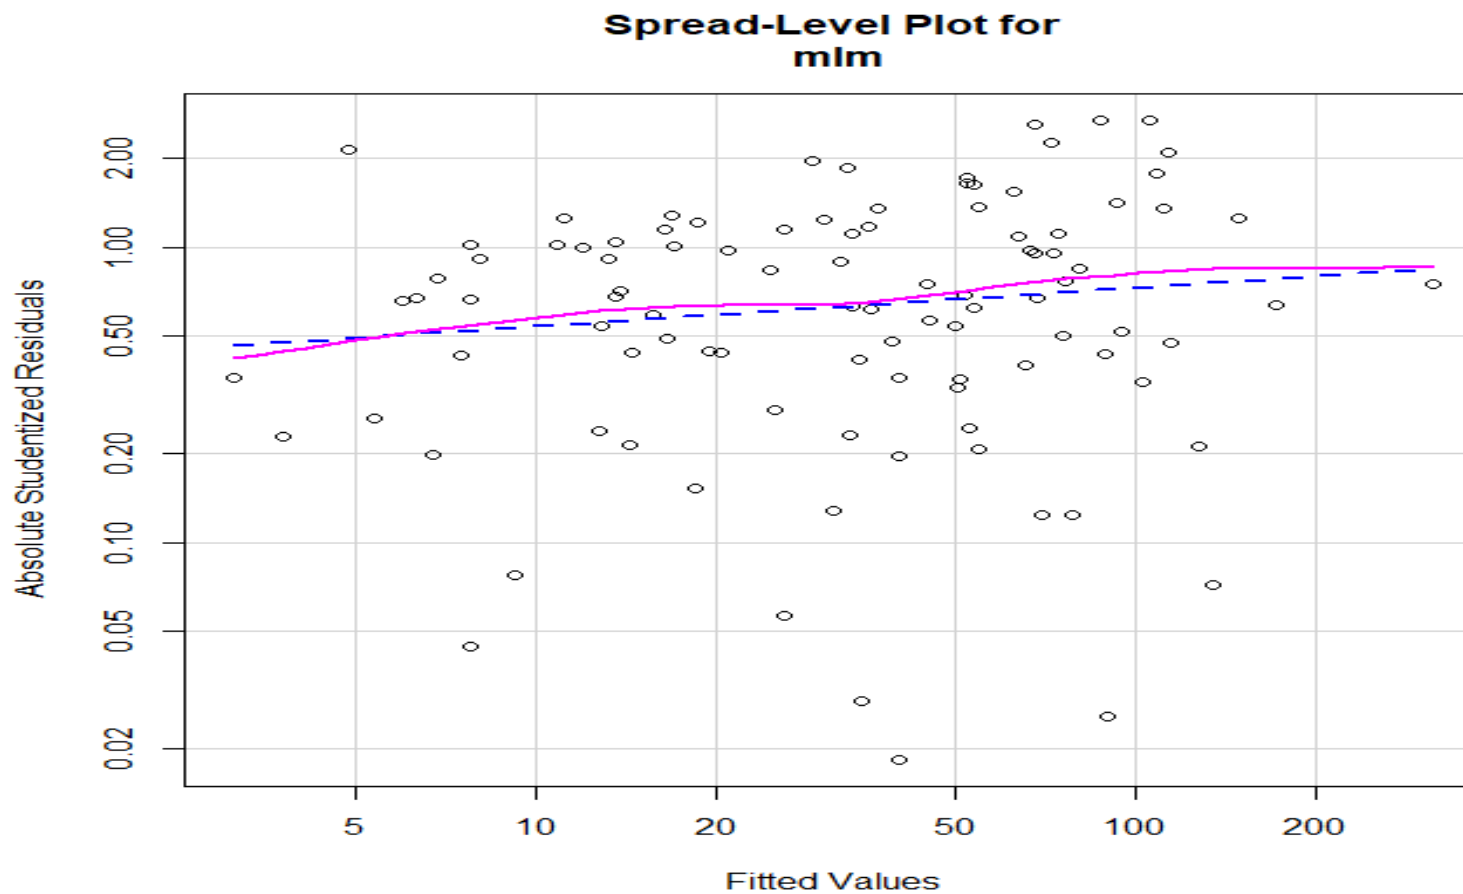
# Parametric Test (normality)

- T-test
- Anova
- Pearson correlation test
- ...

# Which Normality Test Should I Use?

- It is preferable that normality be assessed both visually and through normality tests, of which the **Shapiro-Wilk test** is recommended.
- The KS or Lillie test should be “carefully” used owing to its low power.

# Evaluate homoscedasticity



`spreadLevelPlot(mlm)`



# Test for homoscedasticity

- **ncvTest(mlm)** # Non-constant Variance Score Test
- Computes a score test of the hypothesis of constant error variance against the alternative that the error variance changes with the level of the response (fitted values), or with a linear combination of predictors.

```
> ncvTest(mlm) # non-constant error variance test
Non-constant Variance Score Test
Variance formula: ~ fitted.values
Chisquare = 2.617584, Df = 1, p = 0.10569
```

## Introduce some test homogeneity of variances



- **leveneTest** # package {car}
- **fligner.test**: (nonparametric) k-sample test
- **mood.test**: two-sample test for a difference in scale parameters
- **ansari.test**: two-sample test for a difference in scale parameters. (testing for equal variance for non-normal samples)
- **bartlett.test**: a parametric test of the null that the variances in each of the groups (samples) are the same.
- **var.test**: an F test to compare the variances of two samples from normal populations.
- ...

# Test for independence

- `durbinWatsonTest(mlm)` # test for **autocorrelated** errors

```
> durbinWatsonTest(mlm) # test for autocorrelated errors
lag Autocorrelation D-W Statistic p-value
  1      -0.09174254      2.158922   0.462
Alternative hypothesis: rho != 0
>
> |
```

### *Residuals: dependence diagnosis*



# Multi-collinearity: variance inflation factors

- $\text{vif} > 10$ : multi-collinearity
- Cohen, J, Cohen, P, West, S. G., and Aiken, L. S. 2003. Applied Multiple Regression/Correlation Analysis for the Behavioral Sciences (3rd ed.). Hillsdale, NJ: Lawrence Erlbaum.
- Kutner, Nachtsheim, and Neter, 2004, Applied Linear Regression Models, Fourth Edition, McGraw-Hill Irwin.)

```






.
> vif(mlm)
      x1      x2      x3    x1:x2    x1:x3
3.296571 1.971189 1.703740 3.188911 3.233037
>
> set.seed(108) ; n= 100 ; beta0 = array( c(3, 2, 1, 1), c(1,4) )
>
> x1=rgamma( n, 1, 1/10 ) ; x2= 1.5*x1 + rnorm( n, sd= 0.5 ) ; x3=rexp( n )
> X= cbind( x1, x2, x3 )
> y= beta0[,1] + beta0[,2]*x1 + beta0[,3]*x2 + beta0[,4]*x1*x2 + rnorm( n, sd=1.5 )
>
> mlm2 = lm(y ~ x1 + x2 + x3 + x1:x2 + x1:x3 )
> vif(mlm2)
      x1      x2      x3    x1:x2    x1:x3
732.252036 753.418106  2.266356  7.620066  3.559061
\ |

```

## (Non-) Linearity

- These functions construct component+residual plots, also called partial-residual plots, for linear and generalized linear models.
- ```
df = data.frame(y,x1,x2,x3,x1.x2=x1*x2,  
x1.x3=x1*x3 )  
lm.fit1 = lm( y ~ ., df) ; crPlots(lm.fit1)
```

## Component + Residual Plots





# Global Validation of Linear Model Assumptions

- Pena, EA and Slate, EH (2006). Global validation of linear model assumptions, *J. Am. Stat. Assoc.*, **101**(473):341-354.
- `library(gvlma)`  
`summary(gvlma(mlm))`

```

Call:
lm(formula = y ~ x1 + x2 + x3 + x1:x2 + x1:x3)

Residuals:
    Min       1Q   Median       3Q      Max
-3.7979 -1.0429 -0.0360  0.9585  3.8980

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.770501    0.392648   7.056 2.89e-10 ***
x1           2.044134    0.031253  65.406 < 2e-16 ***
x2           1.075617    0.092715  11.601 < 2e-16 ***
x3          -0.225363    0.229452  -0.982  0.329
x1:x2         0.991471    0.008786 112.846 < 2e-16 ***
x1:x3        -0.003571    0.024526  -0.146  0.885
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.52 on 94 degrees of freedom
Multiple R-squared:  0.9989,    Adjusted R-squared:  0.9989
F-statistic: 1.731e+04 on 5 and 94 DF,  p-value: < 2.2e-16

ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
Level of Significance = 0.05

Call:
gvlma(x = mlm)

              Value p-value              Decision
Global Stat    0.90519  0.9238 Assumptions acceptable.
Skewness       0.37302  0.5414 Assumptions acceptable.
Kurtosis       0.00225  0.9622 Assumptions acceptable.
Link Function  0.49035  0.4838 Assumptions acceptable.
Heteroscedasticity 0.03957  0.8423 Assumptions acceptable.

```

# Comparing models: ANOVA

```
> mlm_maintterms = lm( y~x1+x2+x3)
```

```
> anova(mlm, mlm_maintterms )
```

Analysis of Variance Table

Model 1:  $y \sim x_1 + x_2 + x_3 + x_1:x_2 + x_1:x_3$

Model 2:  $y \sim x_1 + x_2 + x_3$

|   | Res.Df | RSS   | Df | Sum of Sq | F      | Pr(>F)        |
|---|--------|-------|----|-----------|--------|---------------|
| 1 | 94     | 217   |    |           |        |               |
| 2 | 96     | 31772 | -2 | -31555    | 6831.8 | < 2.2e-16 *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

|

```

> mlm = lm(y ~ x1 + x2 + x3 + x1:x2 + x1:x3 )
>
> mlm_correct = lm( y~x1+x2+x1:x2)
>
> anova(mlm, mlm_correct )
Analysis of Variance Table

Model 1: y ~ x1 + x2 + x3 + x1:x2 + x1:x3
Model 2: y ~ x1 + x2 + x1:x2
  Res.Df    RSS Df Sum of Sq    F Pr(>F)
1     94 217.09
2     96 221.49 -2    -4.4042 0.9535 0.3891
\

```

# Model selection

```
require(MASS)
```

```
mlm_all=lm(y~x1+x2+x3+x1:x2+x1:x3+x2:x3 )
```

```
data.input = data.frame( y, x1, x2, x3, x1*x2, x1*x3,  
x2*x3 )
```

```
fit.full = lm(y ~ ., data.input) ; fit.null = lm(y ~ 1,  
data.input)
```

```

Start:  AIC=762.14
y ~ 1

      Df Sum of Sq  RSS  AIC
+ x1...x2  1    182385 17712 521.68
+ x1      1    132044  68054 656.29
+ x1...x3  1     86219 113879 707.77
+ x2      1     21689 178408 752.67
+ x2...x3  1      9107 190990 759.48
<none>          200097 762.14
+ x3      1         27 200070 764.13

Step:  AIC=521.68
y ~ x1...x2

      Df Sum of Sq  RSS  AIC
+ x1      1    17176.8   535.4 173.79
+ x1...x3  1    4219.1 13493.1 496.48
+ x2      1    3624.1 14088.1 500.79
+ x2...x3  1    1664.1 16048.2 513.82
<none>          17712.2 521.68
+ x3      1     336.2 17376.0 521.77

Step:  AIC=173.79
y ~ x1...x2 + x1

      Df Sum of Sq  RSS  AIC
+ x2      1    313.921 221.49  87.521
+ x2...x3  1    128.132 407.28 148.433
<none>          535.41 173.786
+ x1...x3  1      4.181 531.23 175.003
+ x3      1      0.032 535.38 175.780

Step:  AIC=87.52
y ~ x1...x2 + x1 + x2

      Df Sum of Sq  RSS  AIC
<none>          221.49 87.521
+ x3      1     4.3552 217.13 87.535
+ x1...x3  1     2.1763 219.31 88.533
+ x2...x3  1     0.4981 220.99 89.296

Call:
lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

Coefficients:
(Intercept)      x1...x2           x1           x2
      2.571         0.991         2.043         1.062

```

stepAIC(fit.null,direction="forward",scope=list(upper=fit.full,lower=fit.null)) # lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)



```
Start:  AIC=89.28
y ~ x1 + x2 + x3 + x1...x2 + x1...x3 + x2...x3
```

|           | Df | Sum of Sq | RSS     | AIC    |
|-----------|----|-----------|---------|--------|
| - x1...x3 | 1  | 0.0       | 212.3   | 87.28  |
| <none>    |    |           | 212.3   | 89.28  |
| - x2...x3 | 1  | 4.8       | 217.1   | 89.51  |
| - x3      | 1  | 7.0       | 219.3   | 90.53  |
| - x2      | 1  | 60.3      | 272.6   | 112.29 |
| - x1      | 1  | 9404.1    | 9616.4  | 468.61 |
| - x1...x2 | 1  | 28938.4   | 29150.7 | 579.51 |

```
Step:  AIC=87.28
y ~ x1 + x2 + x3 + x1...x2 + x2...x3
```

|           | Df | Sum of Sq | RSS     | AIC    |
|-----------|----|-----------|---------|--------|
| <none>    |    |           | 212.3   | 87.28  |
| - x2...x3 | 1  | 4.8       | 217.1   | 87.53  |
| - x3      | 1  | 8.7       | 221.0   | 89.30  |
| - x2      | 1  | 60.4      | 272.7   | 110.31 |
| - x1      | 1  | 13364.8   | 13577.1 | 501.10 |
| - x1...x2 | 1  | 30734.8   | 30947.1 | 583.49 |

```
Call:
lm(formula = y ~ x1 + x2 + x3 + x1...x2 + x2...x3, data = data.input)
```

Coefficients:

|             |        |        |         |         |
|-------------|--------|--------|---------|---------|
| (Intercept) | x1     | x2     | x3      | x1...x2 |
| 3.3680      | 2.0354 | 0.8701 | -0.6948 | 0.9933  |
| x2...x3     |        |        |         |         |
| 0.1460      |        |        |         |         |

`stepAIC(fit.full,direction="backward")`

```

Start:   AIC=762.14
y ~ 1

      Df Sum of Sq    RSS    AIC
+ x1...x2  1      182385 177112 521.68
+ x1      1      132044  68054 656.29
+ x1...x3  1      86219 113879 707.77
+ x2      1      21689 178408 752.67
+ x2...x3  1        9107 190990 759.48
<none>                 200097 762.14
+ x3      1         27 200070 764.13

Step:   AIC=521.68
y ~ x1...x2

      Df Sum of Sq    RSS    AIC
+ x1      1      17177   535 173.79
+ x1...x3  1       4219 13493 496.48
+ x2      1       3624 14088 500.79
+ x2...x3  1       1664 16048 513.82
<none>                 177112 521.68
+ x3      1        336 17376 521.77
- x1...x2  1      182385 200097 762.14

Step:   AIC=173.79
y ~ x1...x2 + x1

      Df Sum of Sq    RSS    AIC
+ x2      1        314   221  87.52
+ x2...x3  1        128   407 148.43
<none>                 535 173.79
+ x1...x3  1         4   531 175.00
+ x3      1         0   535 175.78
- x1      1      17177 177112 521.68
- x1...x2  1     67518 68054 656.29

Step:   AIC=87.52
y ~ x1...x2 + x1 + x2

      Df Sum of Sq    RSS    AIC
<none>                 221  87.52
+ x3      1         4.4  217  87.53
+ x1...x3  1         2.2  219  88.53
+ x2...x3  1         0.5  221  89.30
- x2      1      313.9   535 173.79
- x1      1     13866.6 14088 500.79
- x1...x2  1    31551.1 31773 582.12

Call:
lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

Coefficients:
(Intercept)      x1...x2           x1           x2
      2.571         0.991         2.043         1.062

```

```

stepAIC(fit.null,direction="both" ,scope=list(upper=fit.full,lower=fit.null
)) # lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

```



```

Start:  AIC=762.14
y ~ 1

      Df Sum of Sq  RSS   AIC
+ x1...x2  1    182385 17712 521.68
+ x1      1    132044  68054 656.29
+ x1...x3  1     86219 113879 707.77
+ x2      1     21689 178408 752.67
+ x2...x3  1       9107 190990 759.48
<none>                 200097 762.14
+ x3      1         27 200070 764.13

Step:  AIC=521.68
y ~ x1...x2

      Df Sum of Sq  RSS   AIC
+ x1      1    17177   535 173.79
+ x1...x3  1     4219 13493 496.48
+ x2      1     3624 14088 500.79
+ x2...x3  1     1664 16048 513.82
<none>                 17712 521.68
+ x3      1      336 17376 521.77
- x1...x2  1    182385 200097 762.14

Step:  AIC=173.79
y ~ x1...x2 + x1

      Df Sum of Sq  RSS   AIC
+ x2      1      314   221  87.52
+ x2...x3  1      128   407 148.43
<none>                 535 173.79
+ x1...x3  1         4   531 175.00
+ x3      1         0   535 175.78
- x1      1    17177 17712 521.68
- x1...x2  1    67518 68054 656.29

Step:  AIC=87.52
y ~ x1...x2 + x1 + x2

      Df Sum of Sq  RSS   AIC
<none>                 221  87.52
+ x3      1         4.4  217  87.53
+ x1...x3  1         2.2  219  88.53
+ x2...x3  1         0.5  221  89.30
- x2      1     313.9   535 173.79
- x1      1    13866.6 14088 500.79
- x1...x2  1    31551.1 31773 582.12

Call:
lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

Coefficients:
(Intercept)          x1...x2              x1              x2
      2.571           0.991           2.043           1.062

```

`stepAIC(fit.null,direction="forward",scope=list(upper=fit.full,lower=fit.null), k=log( n )) # lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)`

```

Start:   AIC=762.14
y ~ 1

      Df Sum of Sq    RSS    AIC
+ x1...x2  1      182385 17712 521.68
+ x1      1      132044  68054 656.29
+ x1...x3  1      86219 113879 707.77
+ x2      1      21689 178408 752.67
+ x2...x3  1        9107 190990 759.48
<none>                 200097 762.14
+ x3      1         27 200070 764.13

Step:   AIC=521.68
y ~ x1...x2

      Df Sum of Sq    RSS    AIC
+ x1      1      17177   535 173.79
+ x1...x3  1       4219 13493 496.48
+ x2      1       3624 14088 500.79
+ x2...x3  1       1664 16048 513.82
<none>                 17712 521.68
+ x3      1        336 17376 521.77
- x1...x2  1      182385 200097 762.14

Step:   AIC=173.79
y ~ x1...x2 + x1

      Df Sum of Sq    RSS    AIC
+ x2      1        314   221  87.52
+ x2...x3  1       128   407 148.43
<none>                 535 173.79
+ x1...x3  1         4   531 175.00
+ x3      1         0   535 175.78
- x1      1      17177 17712 521.68
- x1...x2  1     67518 68054 656.29

Step:   AIC=87.52
y ~ x1...x2 + x1 + x2

      Df Sum of Sq    RSS    AIC
<none>                 221  87.52
+ x3      1         4.4  217  87.53
+ x1...x3  1         2.2  219  88.53
+ x2...x3  1         0.5  221  89.30
- x2      1       313.9   535 173.79
- x1      1     13866.6 14088 500.79
- x1...x2  1    31551.1 31773 582.12

Call:
lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

Coefficients:
(Intercept)      x1...x2           x1           x2
      2.571         0.991         2.043         1.062

```

stepAIC(fit.full,direction="backward", k=log( n ) ) #

lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

```

Start:   AIC=762.14
y ~ 1

      Df Sum of Sq    RSS    AIC
+ x1...x2  1      182385 177112 521.68
+ x1      1      132044  68054 656.29
+ x1...x3  1      86219 113879 707.77
+ x2      1      21689 178408 752.67
+ x2...x3  1        9107 190990 759.48
<none>                 200097 762.14
+ x3      1         27 200070 764.13

Step:   AIC=521.68
y ~ x1...x2

      Df Sum of Sq    RSS    AIC
+ x1      1      17177   535 173.79
+ x1...x3  1       4219 13493 496.48
+ x2      1       3624 14088 500.79
+ x2...x3  1       1664 16048 513.82
<none>                 177112 521.68
+ x3      1        336 17376 521.77
- x1...x2  1      182385 200097 762.14

Step:   AIC=173.79
y ~ x1...x2 + x1

      Df Sum of Sq    RSS    AIC
+ x2      1        314   221  87.52
+ x2...x3  1        128   407 148.43
<none>                 535 173.79
+ x1...x3  1          4   531 175.00
+ x3      1          0   535 175.78
- x1      1      17177 177112 521.68
- x1...x2  1     67518 68054 656.29

Step:   AIC=87.52
y ~ x1...x2 + x1 + x2

      Df Sum of Sq    RSS    AIC
<none>                 221  87.52
+ x3      1         4.4  217  87.53
+ x1...x3  1         2.2  219  88.53
+ x2...x3  1         0.5  221  89.30
- x2      1      313.9   535 173.79
- x1      1     13866.6 14088 500.79
- x1...x2  1    31551.1 31773 582.12

Call:
lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)

Coefficients:
(Intercept)      x1...x2           x1           x2
      2.571         0.991         2.043         1.062

```

`stepAIC(fit.null,direction="both" ,scope=list(upper=fit.full,lower=fit.null`  
`), k=log( n )) # lm(formula = y ~ x1...x2 + x1 + x2, data = data.input)`

## Select on all subsets

```
library(leaps)
data.input = data.frame( y, x1, x2, x3, x1*x2,
x1*x3, x2*x3 )
model.com = regsubsets(y~., nvmax=30,
data=data.input)
model.sum = summary(model.com)
names(model.sum) # "which" "rsq" "rss"
"adjr2" "cp" "bic" "outmat" "obj"
```


|   |       | x1  | x2  | x3  | x1...x2 | x1...x3 | x2...x3 |
|---|-------|-----|-----|-----|---------|---------|---------|
| 1 | ( 1 ) | " " | " " | " " | "*"     | " "     | " "     |
| 2 | ( 1 ) | "*" | " " | " " | "*"     | " "     | " "     |
| 3 | ( 1 ) | "*" | "*" | " " | "*"     | " "     | " "     |
| 4 | ( 1 ) | "*" | "*" | "*" | "*"     | " "     | " "     |
| 5 | ( 1 ) | "*" | "*" | "*" | "*"     | " "     | "*"     |
| 6 | ( 1 ) | "*" | "*" | "*" | "*"     | "*"     | "*"     |

```

which.min(model.sum$rss)    # 6
which.min(model.sum$adjr2) # 1
which.min(model.sum$cp)     # 5
which.min(model.sum$bic)    # 3


```

# Adjusted R-squared




```
plot(model.com,scale="adjr2")
```

# Cp (AIC)



`plot(model.com,scale="Cp")`

# BIC



`plot(model.com,scale="bic")`